

Orchestrating Your Life with AI Agents

A systems approach, not a coding one

Four ways to use AI for ongoing projects

A trip. A health question that lasts a year. A side project. The people you owe replies to.

One question, one chat

New conversation, paste in the context, ask, get an answer.

Next time: explain it again.

Read AI search summaries

Read the AI summary at the top of search results.

60% of U.S. adults read these summaries; 49% use chatbots.

Overlapping groups — Pew, June 2026

Let the app remember

Use the app's memory or project features to keep context.

The context is somewhere. In their app.

Paste or attach context

41% of surveyed U.S. AI users have uploaded a file, image, or document for analysis.

That measures ever-use, not use every session.

Epoch AI/Ipsos, 2026

My test for any approach, three months later: **where does this stand, what did I decide, and why?**

Where it breaks, once a project outlives a conversation

The same four failures, in every tool I tried

You re-explain. Every session starts from the model's memory of you, not the project's record.

You can't update, only append. Things change; the folder just gets one more conversation.

The record isn't yours. It's in a vendor's app, in a vendor's format, on a vendor's timeline.

The "why" is gone. You can find that you kept the hotel and changed the bus. Not why.

And whatever fixes this can't become a job. Keeping the system running has to stay small.

How I got here

Three steps, each ended by something that broke

STEP 1

Projects in the chat apps

A folder of related conversations, some memory across them.

Useful. A place to hold chats.

Broke: I couldn't update the project as things evolved. I just kept adding conversations.

STEP 2

A skill per area of life

Put the context in a skill file. Call it whenever the topic comes up. Let the assistant update it.

Broke: a skill is a standard operating procedure. It was never built to hold a record over months.

STEP 3 — NOW

The record in a folder; the skill, a pointer

One folder of plain files per project: purpose, status, decisions with reasons, next actions, sources, rules for the assistant.

The skill shrinks to a few lines: go here, read these first.

What it costs: a few minutes a session.

Each step was fine until the project outlived it.

Orchestrating your life: the view from 30,000 feet

Not one assistant. One person, one folder of workspaces, any agentic harness, any device.

ONE FOLDER ON YOUR MACHINE · SYNCED TO A PRIVATE BACKUP




Every box is the same thing: a folder of plain files with a purpose, a status, decisions with reasons, and next actions. Same rules, same shape, so any harness knows its way around any of them.

Getting in takes one command, from any harness, on any device: `/trip` at the desk, the same thing from the phone in the car. Assistants and models get swapped; the folders stay.

Who does the work, where the record lives, what can be replaced

■ owned — keep for years □ replaceable — pick per task, per year □ you — the only constant



Try it: one workspace for one real project



`https://github.com/rdmayo21/agentic-workspaces`

Templates + pointer-skill setup for Claude Code, Codex, Gemini CLI. MIT.

Use it for a month. Then tell me what I'm missing.